```
system:
  name: TSAR_RAPTOR_V6
  role: non_blocking_detection_subsystem
  philosophy:
    - decision_over_label
    - uncertainty_is_asset
    - attention_is_governed
    - fail_open_not_fail_close

hardware_envelope:
  cpu:
    max_parallel_jobs: 32
    soft_cap: 24
  gpu:
    max_concurrent_batches: 4
    soft_cap: 2
    vram_soft_limit: 0.85

safety_valve:
  modes:
    - NORMAL
    - DEGRADED_LOW_ONLY
    - BYPASS
```

```
  manual_override:
    enabled: true
    default_mode: NORMAL
    ttl_minutes: 180
    audit_log: true

  auto_circuit_breaker:
    triggers:
      gpu_pressure:
        condition: "gpu.vram_utilization > 0.92"
        action: DEGRADED_LOW_ONLY
      latency_violation:
        condition: "p95_latency_ms > slo.p95_latency_ms"
        action: BYPASS
      human_budget_exhausted:
        condition: "human.today_escalations >= human.max_daily_samples"
        action: DEGRADED_LOW_ONLY
      low_information_period:
        condition: "rolling_1h.mean_information_gain <
policy.minimum_viable_ig"
        action: DEGRADED_LOW_ONLY

slo:
  p95_latency_ms: 300
```

```
    timeout_rate: 0.005
    gpu_budget_share: 0.20
    cpu_budget_share: 0.25
    max_high_cost_rate: 0.05

sla:
    non_blocking: true
    fail_mode: FAIL_OPEN
    enforce_resource_caps: true

scheduler:
    states:

        INIT:
            on_enter:
                - reset_cost
                - reset_uncertainty
            next: LOW_COST

        LOW_COST:
            modules: low_cost
            evaluate:
                uncertainty_remaining: U0_remain
                uncertainty_reduction: U0_reduce
```

```
      information_gain: I0
    decision:
      - if: "I0 < policy.min_continue"
        next: STOP
      - else: MID_COST

MID_COST:
    modules: mid_cost
    evaluate:
      uncertainty_remaining: U1_remain
      uncertainty_reduction: U1_reduce
      disagreement: D1
      information_gain: I1
    decision:
      - if: "I1 < policy.min_continue"
        next: STOP
      - if: "U1_remain < policy.low_exit and D1 < policy.disagreement_high"
        next: ROUTE
      - else: HIGH_COST

HIGH_COST:
    modules: high_cost
    evaluate:
      uncertainty_remaining: U2_remain
```

```
    uncertainty_reduction: U2_reduce
    disagreement: D2
  decision:
    - if: "D2 >= policy.disagreement_high"
      next: HUMAN_ESCALATION
    - else: ROUTE

HUMAN_ESCALATION:
  constraints:
    max_daily_samples: 200
  output:
    - hypothesis_adjustment
    - module_trust_feedback
  terminal: true

ROUTE:
  description: decision_only_non_label
  output:
    - routing_decision
    - confidence_estimate
    - uncertainty_remaining
    - disagreement_score
  terminal: true
```

```
    STOP:
      output:
        - stop_reason
        - info_density_estimate
        - cost_spent
      terminal: true

policy:
  min_continue: 0.10
  low_exit: 0.15
  disagreement_high: 0.35

policy_learning:
  enabled: true
  objective:
    maximize:
      - uncertainty_reduction_per_cost
    minimize:
      - wasted_high_cost_calls
  adjustable_thresholds:
    min_continue: [0.05, 0.25]
    disagreement_high: [0.25, 0.50]
  safety:
    max_delta_per_update: 0.05
```

```
        rollback_on_regression: true


attention_budget_governance:
    core_integrity:
        floor_gpu: 1
        priority: 10
    detection_layer:
        ceiling_gpu: 1
        throttleable: true
        priority: 7
    economic_layer:
        ceiling_gpu: 2
        priority: 5
    degradation_order:
        - economic_layer
        - detection_layer
        - core_integrity
```